



zenika
<animés par la passion>

Elastic Suite Configuration Details

Version 1.18-SNAPSHOT

2019-01-22

Table of Contents

1. Elastic servers configuration	2
1.1. Initialize VM	2
1.2. Install Docker	3
1.2.1. Define Nexus3 as the Docker registry	4
1.3. Setup a dockerized Oracle12c database	6
1.4. Install Elastic items	7
1.4.1. Migration prerequisites	7
1.4.2. Elasticsearch	8
1.5. Kibana	9
1.5.1. Troubleshoot	11
1.6. Curator	13
1.6.1. Automation	13
1.6.2. Configuration	13
1.7. Heartbeat	15
1.8. Logstash	17
1.9. Filebeat	18
1.10. Metricbeat	19
1.11. Grafana	21
1.12. Jaeger Tracing (OpenZipkin-like)	22
2. Appendix	23
2.1. Revision marks	23

Table 1. History

Date	Author	Detail
2019-01-22	bcouetil	plugins slides + fragmented examples + text wrapping near images + screenshot-as-code + big-image + charts axis labels + css fixes
2018-12-18	bcouetil	Moved Gerrit plugin configuration from Jenkins to Gerrit page + Enhanced HTML CSS + Cropped some images + Improved PlantUML skin
2018-08-29	bcouetil	Asciidoc HTML look & feel changes
2018-08-24	bcouetil	Icones added for download + favicon added for webpage
2018-08-23	bcouetil	Initial commit

1. Elastic servers configuration

1.1. Initialize VM

- Adding a user

```
$ adduser devops
```

- Granting him root privileges

```
$ visudo
```

```
devops ALL=(ALL:ALL) ALL
```

- Checking FS size

```
$ parted
(parted) $ print free

Number  Start   End     Size    Type    File system  Flags
1       32.3kB  1049kB  1016kB             Free Space
1       1049kB  85.9GB  85.9GB  primary xfs          boot
        85.9GB  85.9GB  16.9kB             Free Space

(parted) $ quit
```

- Example

Number	Start	End	Size	Type	File system	Flags
	32.3kB	1049kB	1016kB		Free Space	
1	1049kB	500MB	499MB	primary	ext2	boot
2	500MB	53.7GB	53.2GB	primary		lvm
	53.7GB	53.7GB	1049kB		Free Space	



Below instructions are for Ubuntu only. You can check your Linux distribution with this command : `cat /etc/*-release`

- Add some server for apt-get

```
$ sudo vi /etc/apt/sources.list
```

```
deb [arch=amd64] http://archive.ubuntu.com/ubuntu/ trusty main restricted universe multiverse
deb [arch=amd64] http://archive.ubuntu.com/ubuntu/ trusty-security main restricted universe multiverse
deb [arch=amd64] http://archive.ubuntu.com/ubuntu/ trusty-updates main restricted universe multiverse
deb [arch=amd64] http://archive.ubuntu.com/ubuntu/ trusty-proposed main restricted universe multiverse
deb [arch=amd64] http://archive.ubuntu.com/ubuntu/ trusty-backports main restricted universe multiverse
```

1.2. Install Docker



Below instructions are for Ubuntu 14 only. You can check your Linux distribution with this command : `cat /etc/*-release`

```
$ apt-get install apt-transport-https ca-certificates curl software-properties-common
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
$ add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu xenial stable"
$ apt-get update
$ apt-cache search docker-ce
$ apt-get install docker-ce
```

- May work on Jenkins slave



don't use on managed PL, we don't have enough rights

```
$ sudo add-apt-repository \
    "deb [arch=amd64] https://download.docker.com/linux/$(. /etc/os-release; echo "$ID") \
    $(lsb_release -cs) \
    stable"
$ sudo apt-get update
$ sudo apt-cache search docker-ce
$ sudo apt-get install --assume-yes docker-ce
$ sudo dockerd
```

- Allow Docker remote API

Solution found here <https://forums.docker.com/t/enable-remote-api-on-docker-hosts-in-ubuntu-14/11583/2>

```
$ vi /etc/default/docker
```

```
$ DOCKER_OPTS="-H tcp://0.0.0.0:2375 -H unix:///var/run/docker.sock"
```

- Start/restart Docker Daemon

Ubuntu

```
$ sudo dockerd
```

```
$ service docker restart
```

centos

```
$ systemctl enable docker
$ systemctl start docker
```

- To restart (as root)



Don't forget the `docker.sock` `chmod` if you use `metricbeat`

```
$ service docker restart
```

- To check FS size

```
root@frpardge:/var/lib/docker
$ du -sh -- * .*
92K   aufs
44K   containers
116K  image
52K   network
20K   plugins
4.0K  swarm
4.0K  tmp
4.0K  trust
28K   volumes
4.0K  .
61M   ..
```

- Get rid of sudo for devops user

```
$ sudo groupadd docker
$ sudo gpasswd -a devops docker
$ newgrp docker
$ docker run hello-world
```

- Install **Portainer** to ease administration

```
$ sudo docker pull portainer/portainer

$ sudo docker run -d --name portainer --restart=always -p 19000:9000 -v /var/run/docker.sock:/var/run/docker.sock
portainer/portainer
```

- To use, go to <http://frpardge.corp.nvx.com:19000>
 - login/password = admin / **
 - Install docker-compose

```
curl -L https://github.com/docker/compose/releases/download/1.19.0/docker-compose-`uname -s`-`uname -m` -o
/usr/local/bin/docker-compose

chmod +x /usr/local/bin/docker-compose

docker-compose --version
```

1.2.1. Define Nexus3 as the Docker registry

- Raise a ticket in INSERE to ask a port opening for Nexus3 as a Docker registry
 - They will provide this kind of response, which indicates how to login before 'docker push' :

```
$ docker login docker-registry-bpmfactory.s2-eu.nvx.com
User name: docker
User Password: dockerPWdbpmfactory
```

- Use the information to add the registry in docker configuration

```
$ vi /etc/docker/daemon.json
```

```
{  
  "storage-driver": "devicemapper",  
  "insecure-registries": [  
    "docker-registry-bpmfactory.s2-eu.nvx.com"  
  ],  
  "disable-legacy-registry": true  
}
```

- be careful not to have INSECURE_REGISTRY here, it would not start :

```
$ vim /etc/sysconfig/docker
```

```
#INSECURE_REGISTRY='--insecure-registry userbxy05.socle:8444'
```

- Redémarrer docker

```
$ service docker restart
```

1.3. Setup a dockerized Oracle12c database

Database found here : <https://hub.docker.com/r/sath89/oracle-12c/>

```
$ docker pull sath89/oracle-12c

$ docker run --restart=always --name dbdev -d -p 18080:8080 -p 1521:1521 sath89/oracle-12c

$ docker logs -f feef20144fdc124d7b19d22aaf7bd63cbb837df667cc9764e7bdb5bcafa1af46

Database not initialized. Initializing database.
Starting tnslnr
Copying database files
1% complete
3% complete
Import finished
Database ready to use. Enjoy! ;)
```

Connect to Oracle Application Express web management console with following settings :

- host = <http://frpardge:18080/apex>
- workspace = **INTERNAL**
- user = **ADMIN**
- password **Oracle!**

1.4. Install Elastic items

Configuration files are given in next associated sections below. For some of them, some chmod change is needed :

```
$ cd ~/elastic
$ chmod go-w ./*.yaml
```

1.4.1. Migration prerequisites

If you are upgrading from a previous version of Elastic, you have to do this before anything :

- Close data senders using Portainer for containers
 - Shutdown the IS, or just disable CgElastic & WmMediator packages
 - Stop Heartbeat, Filebeat, Metricbeat containers
 - No need to stop Logstash if Filebeat is closed
- Check that nothing is coming in Elasticsearch with Kibana, then stop Kibana container
- Stop Elasticsearch container

For now, no data migration has been tried, so no support on it. This will be a fresh new Elasticsearch, and a Kibana with imported dashboards (hoping they still work).

Rename all stopped container, to be able to get the initial name on new containers.

1.4.2. Elasticsearch



If you are new to the Elastic Stack, learn with the excellent official Kibana tutorial : <https://www.elastic.co/guide/en/kibana/current/getting-started.html>

- Install with docker without x-pack

```
$ docker pull docker.elastic.co/elasticsearch/elasticsearch-oss:6.0.0
```

To start it

```
$ docker run --restart=always -d --name elastic -p 9200:9200 -p 9300:9300 -e "discovery.type=single-node"
docker.elastic.co/elasticsearch/elasticsearch-oss:6.0.0
```

- if elastic stops directly after start with this error

```
max virtual memory areas vm.max_map_count [65530] likely too low, increase to at least [262144]
```

- Then type before retry

```
$ sudo sysctl -w vm.max_map_count=262144
```

1.5. Kibana

- Install with docker without x-pack

```
$ docker pull docker.elastic.co/kibana/kibana-oss:6.0.0
```

- Create the file described at the end of this section

```
~/elastic/kibana.yml
```

- Start the container

```
$ docker run --restart=always -d --name kibana -p 5601:5601 -v ~/elastic/kibana.yml:/usr/share/kibana/config/kibana.yml
docker.elastic.co/kibana/kibana-oss:6.0.0
```

Check that it is up and running : <http://frpardge:5601/>

Once every application is up, you will be able to declare patterns :

- cgwmbeat-*
- heartbeat-*
- jenkins
- logstash-*
- metricbeat-*
- webmethodsmediator

And to apply some Elasticsearch default index configuration :

- the limit of 1000 fields by index is a bit low, updated to 2000
- default is 5 shards per index, too many for dev
- default is 1 replica, for a single node ES it's 0

```
PUT _template/all
{
  "index_patterns" : ["*"],
  "settings": {
    "index.mapping.total_fields.limit": 2000,
    "index.max_docvalue_fields_search": 400,
    "number_of_shards": 1,
    "number_of_replicas": 0
  }
}
```

Here is something to try, inside the "PUT _template/all", someday, to not have keyword (fixed word) + text (searchable) but only keyword :

```
"dynamic_templates": [
  {
    "match_mapping_type": "string",
    "mapping": {
      "type": "keyword"
    }
  }
]
```

For Elasticsearch monitoring :

```
GET /_cat/indices?v
GET _cluster/health
```

~/elastic/kibana.yml

```
# Kibana is served by a back end server. This setting specifies the port to use.
#server.port: 5601

# Specifies the address to which the Kibana server will bind. IP addresses and host names are both valid values.
# The default is 'localhost', which usually means remote machines will not be able to connect.
# To allow connections from remote users, set this parameter to a non-loopback address.
server.host: "0.0.0.0"

# Enables you to specify a path to mount Kibana at if you are running behind a proxy. This only affects
# the URLs generated by Kibana, your proxy is expected to remove the basePath value before forwarding requests
# to Kibana. This setting cannot end in a slash.
#server.basePath: ""

# The maximum payload size in bytes for incoming server requests.
#server.maxPayloadBytes: 1048576

# The Kibana server's name. This is used for display purposes.
#server.name: "your-hostname"

# The URL of the Elasticsearch instance to use for all your queries.
elasticsearch.url: "http://frpardge.corp.nvx.com:9200"

# When this setting's value is true Kibana uses the hostname specified in the server.host
# setting. When the value of this setting is false, Kibana uses the hostname of the host
# that connects to this Kibana instance.
#elasticsearch.preserveHost: true

# Kibana uses an index in Elasticsearch to store saved searches, visualizations and
# dashboards. Kibana creates a new index if the index doesn't already exist.
#kibana.index: ".kibana"

# The default application to load.
#kibana.defaultAppId: "discover"

# If your Elasticsearch is protected with basic authentication, these settings provide
# the username and password that the Kibana server uses to perform maintenance on the Kibana
# index at startup. Your Kibana users still need to authenticate with Elasticsearch, which
# is proxied through the Kibana server.
#elasticsearch.username: "user"
#elasticsearch.password: "pass"

# Enables SSL and paths to the PEM-format SSL certificate and SSL key files, respectively.
# These settings enable SSL for outgoing requests from the Kibana server to the browser.
#server.ssl.enabled: false
#server.ssl.certificate: /path/to/your/server.crt
#server.ssl.key: /path/to/your/server.key

# Optional settings that provide the paths to the PEM-format SSL certificate and key files.
# These files validate that your Elasticsearch backend uses the same key files.
#elasticsearch.ssl.certificate: /path/to/your/client.crt
#elasticsearch.ssl.key: /path/to/your/client.key
```

```
# Optional setting that enables you to specify a path to the PEM file for the certificate
# authority for your Elasticsearch instance.
#elasticsearch.ssl.certificateAuthorities: [ "/path/to/your/CA.pem" ]

# To disregard the validity of SSL certificates, change this setting's value to 'none'.
#elasticsearch.ssl.verifyMode: full

# Time in milliseconds to wait for Elasticsearch to respond to pings. Defaults to the value of
# the elasticsearch.requestTimeout setting.
#elasticsearch.pingTimeout: 1500

# Time in milliseconds to wait for responses from the back end or Elasticsearch. This value
# must be a positive integer.
#elasticsearch.requestTimeout: 30000

# List of Kibana client-side headers to send to Elasticsearch. To send *no* client-side
# headers, set this value to [] (an empty list).
#elasticsearch.requestHeadersWhitelist: [ authorization ]

# Header names and values that are sent to Elasticsearch. Any custom headers cannot be overwritten
# by client-side headers, regardless of the elasticsearch.requestHeadersWhitelist configuration.
#elasticsearch.customHeaders: {}

# Time in milliseconds for Elasticsearch to wait for responses from shards. Set to 0 to disable.
#elasticsearch.shardTimeout: 0

# Time in milliseconds to wait for Elasticsearch at Kibana startup before retrying.
#elasticsearch.startupTimeout: 5000

# Specifies the path where Kibana creates the process ID file.
#pid.file: /var/run/kibana.pid

# Enables you specify a file where Kibana stores log output.
#logging.dest: stdout

# Set the value of this setting to true to suppress all logging output.
#logging.silent: false

# Set the value of this setting to true to suppress all logging output other than error messages.
#logging.quiet: false

# Set the value of this setting to true to log all events, including system usage information
# and all requests.
#logging.verbose: false

# Set the interval in milliseconds to sample system and process performance
# metrics. Minimum is 100ms. Defaults to 5000.
#ops.interval: 5000
```

1.5.1. Troubleshoot

Here is a list of problems and solutions.

Kibana cannot connect to Elasticsearch

If Kibana cannot connect to Elasticsearch with this message :

```
blocked by: [FORBIDDEN/12/index read-only / allow delete (api)]; [cluster_block_exception] blocked by:
[FORBIDDEN/12/index read-only / allow delete (api)];
```

Then apply these settings :

```
PUT _settings
{
  "index": {
    "blocks": {
      "read_only_allow_delete": "false"
    }
  }
}

PUT cgwmbeat-2018.02.16/_settings
{
  "index": {
    "blocks": {
      "read_only_allow_delete": "false"
    }
  }
}
```

1.6. Curator

```
$ wget -qO - https://packages.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -  
$ sudo vi /etc/apt/sources.list  
deb [arch=amd64] http://packages.elastic.co/curator/5/debian stable main  
$ sudo apt-get update && sudo apt-get install elasticsearch-curator
```

To start it

```
$ curator --config ~/elastic/curator.config.yml --dry-run ~/elastic/curator.delete_indices.yml  
$ curator --config ~/elastic/curator.config.yml ~/elastic/curator.delete_indices.yml
```

1.6.1. Automation

- Create below script

~/elastic/curator.sh

```
#!/bin/sh  
curator --config ~/elastic/curator.config.yml ~/elastic/curator.delete_indices.yml
```

- Open crontab

```
$ crontab -e
```

- Add this line to launch it at 8:00 everyday

```
0 8 * * * ~/elastic/curator.sh
```

- Exit and save with **Ctrl+X**, **Y**, **Enter**

1.6.2. Configuration

~/elastic/curator.config.yml

```
---
# Remember, leave a key empty if there is no value.  None will be a string,
# not a Python "NoneType"
client:
  hosts:
    - 127.0.0.1
  port: 9200
  url_prefix:
  use_ssl: False
  certificate:
  client_cert:
  client_key:
  ssl_no_validate: False
  http_auth:
  timeout: 30
  master_only: False

logging:
  loglevel: INFO
  logfile:
  logformat: default
  blacklist: ['elasticsearch', 'urllib3']
```

~/elastic/curator.delete_indices.yml

```
---
# Remember, leave a key empty if there is no value.  None will be a string,
# not a Python "NoneType"
#
# Also remember that all examples have 'disable_action' set to True.  If you
# want to use this action as a template, be sure to set this to False after
# copying it.
# # # #
# curator --config ~/elastic/curator.config.yml --dry-run ~/elastic/curator.delete_indices.yml
# curator --config ~/elastic/curator.config.yml ~/elastic/curator.delete_indices.yml
# # # #
actions:
  1:
    action: delete_indices
    description: Delete indices older than 30 days. No error when no actual deletion.
    options:
      ignore_empty_list: True
    filters:
      - filtertype: age
        source: name
        direction: older
        timestring: '%Y.%m.%d'
        unit: days
        unit_count: 30
```


1.7. Heartbeat

- Pull the image

```
$ docker pull docker.elastic.co/beats/heartbeat:6.0.0
```

- Create the file described at the end of this section

```
~/elastic/heartbeat.yml
```

- Start the container

```
$ docker run --name heartbeat -d -v ~/elastic/heartbeat.yml:/usr/share/heartbeat/heartbeat.yml  
docker.elastic.co/beats/heartbeat:6.0.0
```

~/elastic/heartbeat.yml

```
#
# wget --user=svc-fr-pldouane --password=Na9Is4Aw0! https://cdsdouane.pl.s2-eu.nvx.com/jenkins/job/DTXE_P1_CodeReview/
#
heartbeat.monitors:
- name: Jenkins
  type: http
  schedule: '@every 30s'
  urls: ["https://bpmfactory.s2-eu.nvx.com/jenkins/job/CNAV-DGE_P1_Review/"]
  username: svc-fr-bpmfact
  password: ****
  check.request.method: GET
  check.response.status: 200
- name: 'Jenkins Douane'
  type: http
  schedule: '@every 30s'
  urls: ["https://cdsdouane.pl.s2-eu.nvx.com/jenkins/job/DTXE_P1_CodeReview/"]
  username: svc-fr-pldouane
  password: ****
  check.request.method: GET
  check.response.status: 200
- name: 'Gerrit home'
  type: http
  schedule: '@every 30s'
  urls: ["https://bpmfactory.s2-eu.nvx.com/gerrit/changes/?n=25&0=81"]
  username: svc-fr-bpmfact
  password: ****
  check.response.status: 200
- name: 'Gerrit viewFile'
  type: http
  schedule: '@every 30s'
  urls: ["https://bpmfactory.s2-
eu.nvx.com/gerrit/changes/421/revisions/5ab9d4c5cab6a087b936748f2df6550666a502dd/files/Jenkinsfile-2-deploy-to-
dev/diff?context=ALL"]
  username: svc-fr-bpmfact
  password: ****
  check.response.status: 200
- name: 'IS Dev'
  type: http
  schedule: '@every 30s'
  urls: ["http://frpardge:5555"]
  username: Administrator
  password: ****
  check.response.status: 200
- name: Kibana
  type: http
  schedule: '@every 30s'
  urls: ["http://frpardge:5601/app/kibana#/management?_g=()"]
  check.response.status: 200
- name: 'UM Dev'
  type: tcp
  schedule: '@every 30s'
  hosts: ["frpardge:9000"]

heartbeat.scheduler:
  limit: 10

output.elasticsearch:
  hosts: ["frpardge.corp.nvx.com:9200"]

dashboards.enabled: true
```

1.8. Logstash



Install this only if you have files to be parsed and sent to Elasticsearch

- Pull the image

```
$ docker pull docker.elastic.co/logstash/logstash-oss:6.0.0
```

- Create the file described at the end of this section

```
~/elastic/logstash-pipelines/logstash.conf
```

- Start the container

```
$ docker run --restart=always --name logstash -d -p 5043:5043 -v ~/elastic/logstash-pipelines/:/usr/share/logstash/pipeline/ docker.elastic.co/logstash/logstash-oss:6.0.0
```

~/elastic/logstash-pipelines/logstash.conf

```
input {
  beats {
    port => "5043"
  }
}

filter {
  if [fields][log_type] == "perflog" {
    grok {
      match => { "message" => "%{TIMESTAMP_ISO8601:timestamp} INFO  PERFORMANCES - \[%{GREEDYDATA:package}\]
%{WORD:method}\\(\\) completed successfully in %{NUMBER:duration:int} ms" }
    }
  }
  else {
    grok {
      match => { "message" => "\[%{TIMESTAMP_ISO8601:timestamp}\\] \[%{NOTSPACE:wMCode}\\] %{GREEDYDATA:textMsg}" }
    }
  }
  date {
    match => [ "timestamp", ISO8601 ]
    timezone => "Europe/Paris"
    target => "@timestamp"
  }
}

output {
  elasticsearch {
    hosts => [ "frpardge.corp.nvx.com:9200" ]
  }
  #stdout { codec => rubydebug }
}
```

1.9. Filebeat



Install this only if you have files to be parsed and sent to Elasticsearch

- Pull the image

```
$ docker pull docker.elastic.co/beats/filebeat:6.0.0
```

- Create the file described at the end of this section

```
~/elastic/filebeat.yml
```

- Start the container

```
$ docker run --name filebeat -d -v /opt/sagis/IntegrationServer/instances/default/logs:/islogs/ -v  
~/elastic/filebeat.yml:/usr/share/filebeat/filebeat.yml docker.elastic.co/beats/filebeat:6.0.0
```

~/elastic/filebeat.yml

```
filebeat.prospectors:  
  
- type: log  
  paths:  
    - /islogs/log4j2/perfs.log  
  fields: {log_type: perflog}  
  
- type: log  
  paths:  
    - /islogs/server.log  
  multiline.pattern: '^\[20'  
  multiline.negate: true  
  multiline.match: after  
  fields: {log_type: serverlog}  
  
output.logstash:  
  hosts: ["frpardge.corp.nvx.com:5043"]
```

1.10. Metricbeat

This chmod has to be done again after each VM reboot before starting Metricbeat :

```
$ sudo chmod 777 /var/run/docker.sock
```

- Pull the image

```
$ docker pull docker.elastic.co/beats/metricbeat:6.0.0
```

- Create the file described at the end of this section

```
~/elastic/metricbeat.yml
```

- Start the container

```
$ docker run --name metricbeat -d -v /var/run/docker.sock:/var/run/docker.sock -v  
~/elastic/metricbeat.yml:/usr/share/metricbeat/metricbeat.yml --volume=/proc:/hostfs/proc:ro --volume  
=/sys/fs/cgroup:/hostfs/sys/fs/cgroup:ro --volume=:/hostfs:ro --net=host docker.elastic.co/beats/metricbeat:6.0.0  
metricbeat -e -system.hostfs=/hostfs
```

To test you CPU graphs, with the proper handling of the cores, you can use stress application to load one or multiple cores :

```
$ sudo apt-get install stress  
$ stress --cpu 2
```

```
metricbeat.modules:
- module: system
  period: 10s
  metricsets:
    - cpu
    #- load
    - memory
    #- network
    - process
    - process_summary
    #- core
    #- diskio
    #- socket
  processes: ['.*']
  process.include_top_n:
    by_cpu: 10      # include top processes by CPU
    by_memory: 10   # include top processes by memory

- module: system
  period: 1m
  metricsets:
    - filesystem
    - fsstat
  processors:
    - drop_event.when.regexp:
        system.filesystem.mount_point: '^/(sys|cgroup|proc|dev|etc|hostfs|run|var)($|/)'

- module: docker
  metricsets:
    #- container
    - cpu
    #- diskio
    #- healthcheck
    #- image
    #- info
    - memory
    #- network
  hosts: ["unix:///var/run/docker.sock"]
  period: 10s

output.elasticsearch:
  hosts: ["frpardge.corp.nvx.com:9200"]

metricbeat.config.modules:
  path: /usr/share/metricbeat/metricbeat.yml
  reload.enabled: true
  reload.period: 60s
```

1.11. Grafana

```
$ wget https://s3-us-west-2.amazonaws.com/grafana-releases/release/grafana_4.4.3_amd64.deb  
$ sudo apt-get install -y adduser libfontconfig  
$ sudo dpkg -i grafana_4.4.3_amd64.deb
```

To start it

```
$ sudo service grafana-server start
```

To auto start it at boot time

```
$ sudo update-rc.d grafana-server defaults
```

1.12. Jaeger Tracing (OpenZipkin-like)

To start it

```
$ docker run --restart=always --name jaeger -d -p5775:5775/udp -p6831:6831/udp -p5778:5778 -p16686:16686  
jaegertracing/all-in-one:latest
```



2. Appendix

2.1. Revision marks

Differences since last tag

